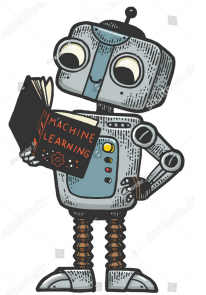


# Introduction to Supervised Learning

MAIN educational 2021

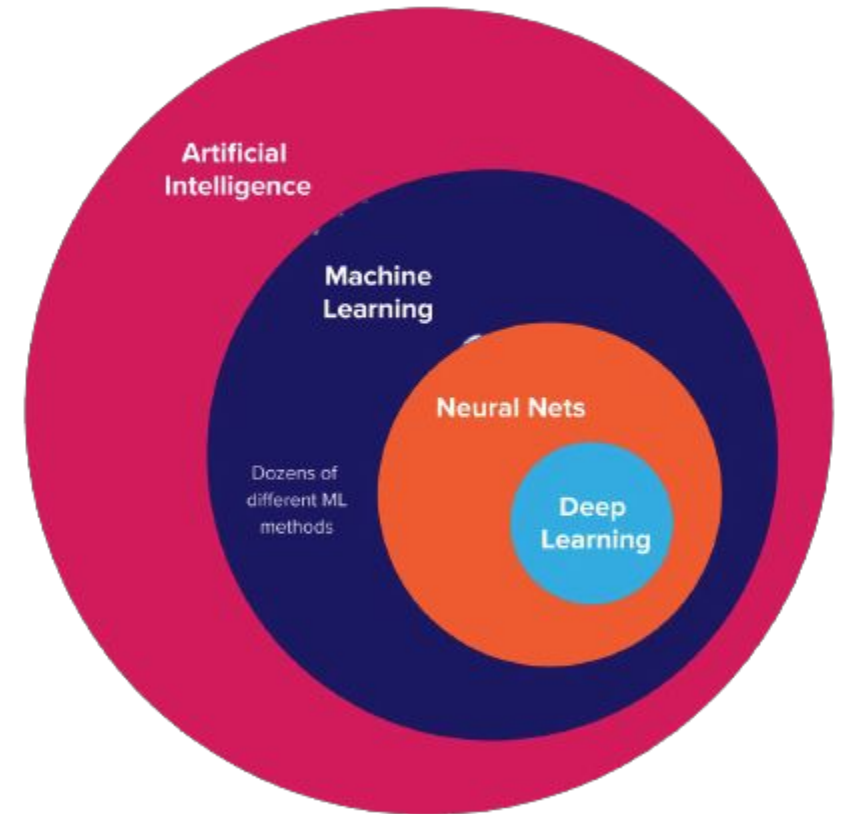
*By*  
*Nikhil Bhagwat & Jérôme Dockès*

25 November 2021



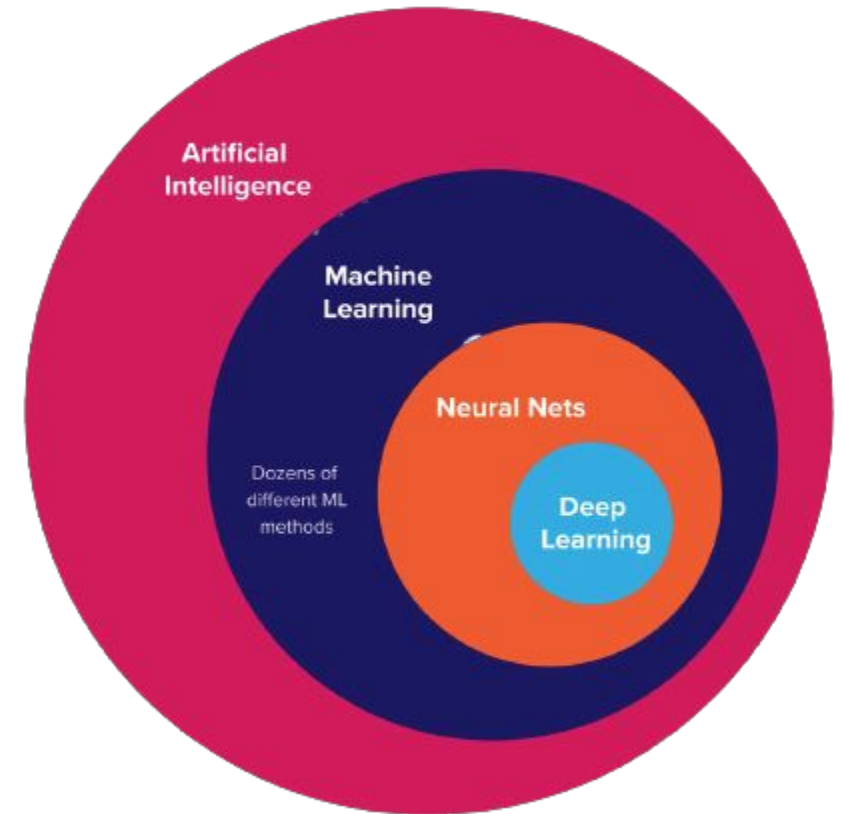
# The Basic Questions

- What is Machine learning (ML)?
  - ML is the study of computer algorithms that improve automatically through experience and by the use of data.




# The Basic Questions

- What is Machine learning (ML)?
  - ML is the study of computer algorithms that improve automatically through experience and by the use of data.
- Why is it useful - especially in life sciences?
  - Biology, Medicine, Environmental sciences comprise phenomena (e.g. a disease) with large number of variables.
  - We want to model complex relationships within these variables.
  - ML can help in these cases and provide accurate predictions.

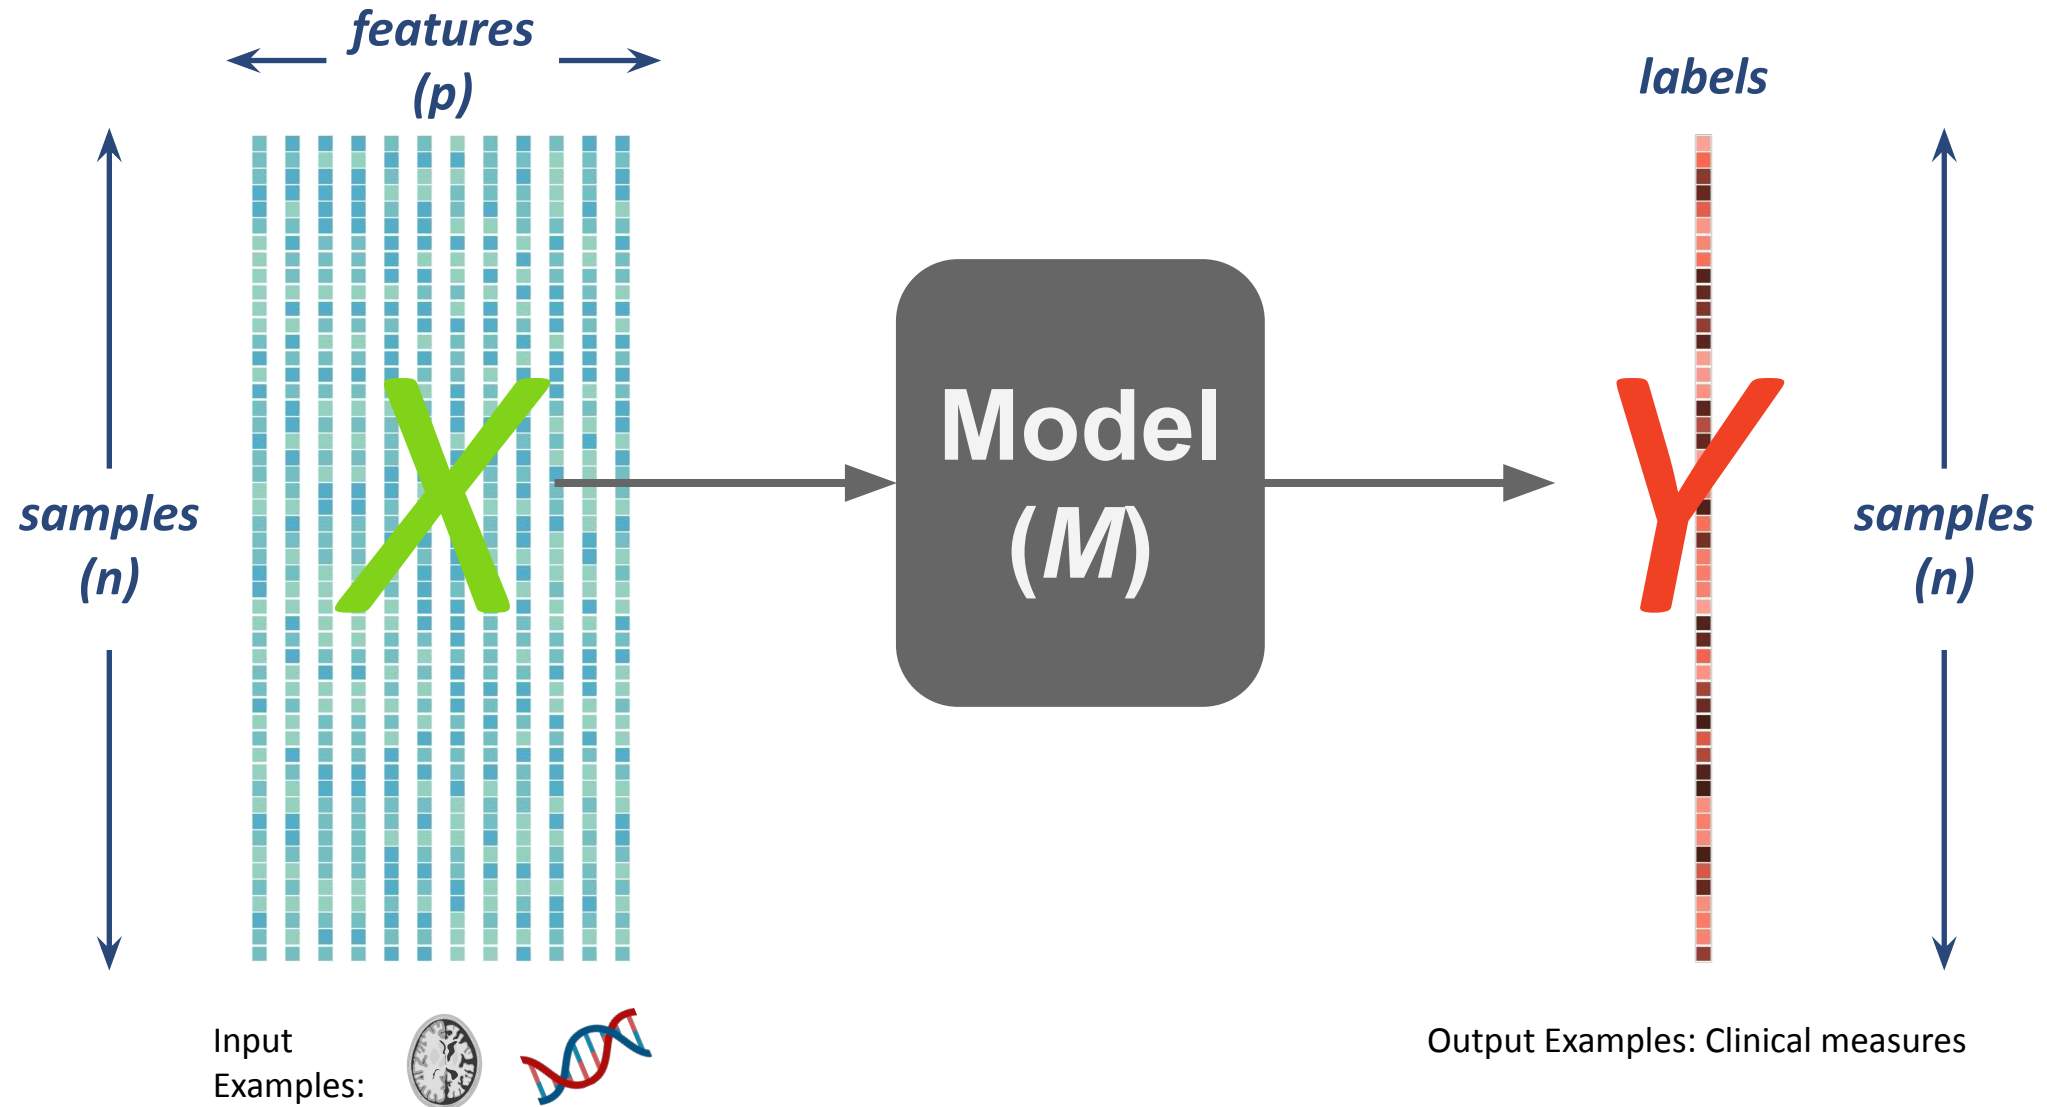


# The Basic Questions

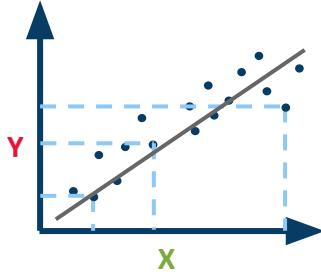
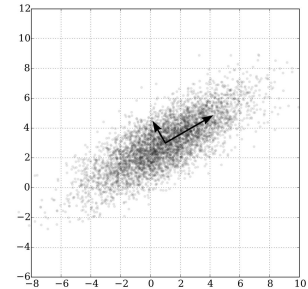
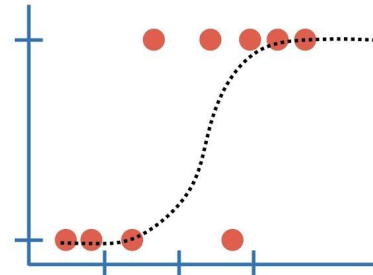
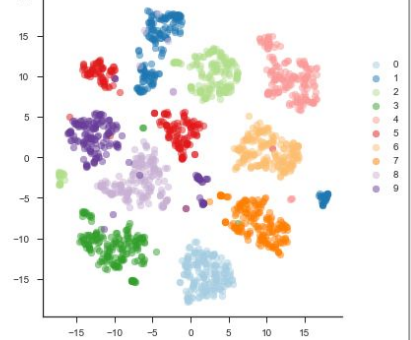
- What is Machine learning (ML)?
  - ML is the study of computer algorithms that improve automatically through experience and by the use of data.
- Why is it useful - especially in life sciences?
  - Biology, Medicine, Environmental sciences comprise phenomena (e.g. a disease) with large number of variables.
  - We want to model complex relationships within these variables.
  - ML can help in these cases and provide accurate predictions.
- When do I use it?
  - You are interested in 1) prediction tasks or 2) low-dimensional representation.
  - You have sufficient data.





# Terminology



# Types of ML Algorithms


Outcome	Supervised Learning	Unsupervised Learning
Continuous	<p>Regression</p> 	<p>Dimensionality reduction</p> 
Categorical	<p>Classification</p> 	<p>Clustering</p> 

# Supervised Learning

- Goal: *Learn* parameters (or weights) of a model (M) that maps  to 

# Supervised Learning $x \rightarrow M \rightarrow y$

- Goal: *Learn* parameters (or weights) of a model ( $M$ ) that maps  $x$  to  $y$
- Example models:
  - Linear / Logistic regression




Linear Regression

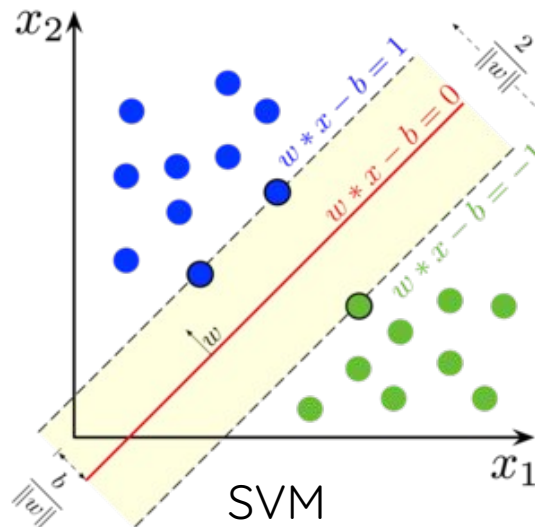


# Supervised Learning $x \rightarrow M \rightarrow y$

- Goal: *Learn* parameters (or weights) of a model ( $M$ ) that maps  $x$  to  $y$
- Example models:
  - Linear / Logistic regression
  - Support vector machines



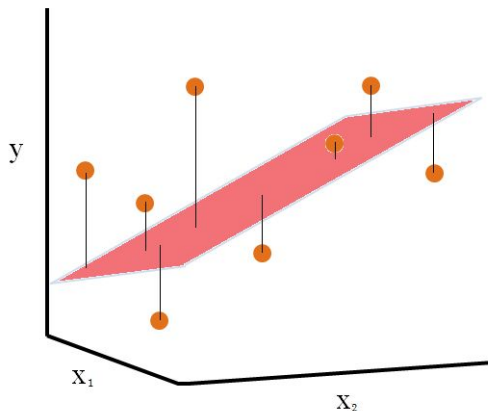
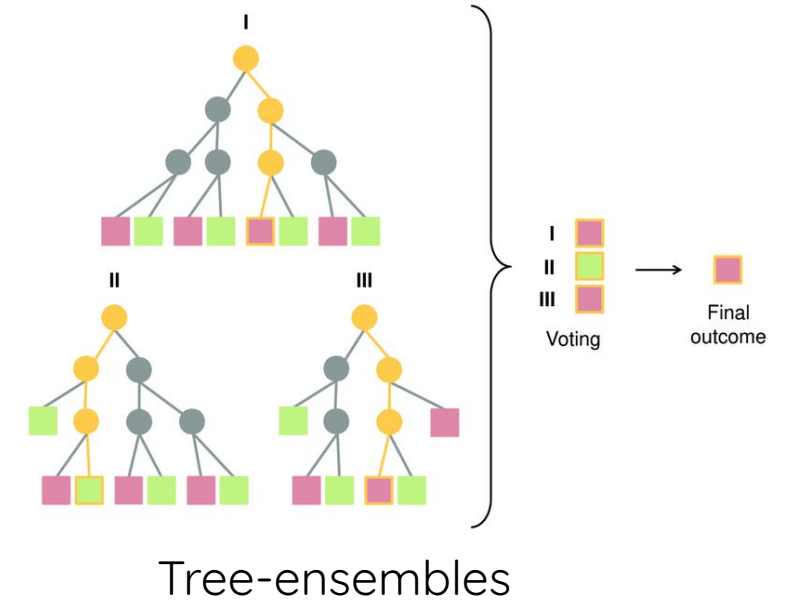
Linear Regression




SVM

# Supervised Learning $x \rightarrow M \rightarrow y$

- Goal: *Learn* parameters (or weights) of a model ( $M$ ) that maps  $x$  to  $y$
- Example models:
  - Linear / Logistic regression
  - Support vector machines
  - Tree-ensembles: random forests, gradient boosting



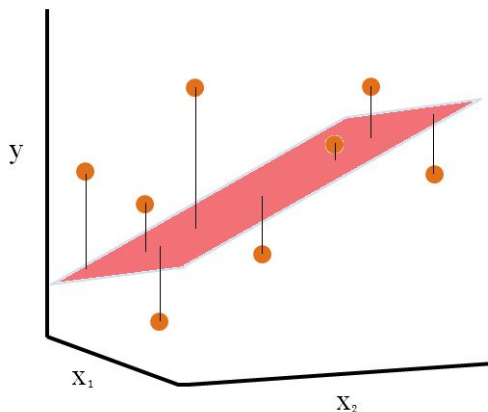
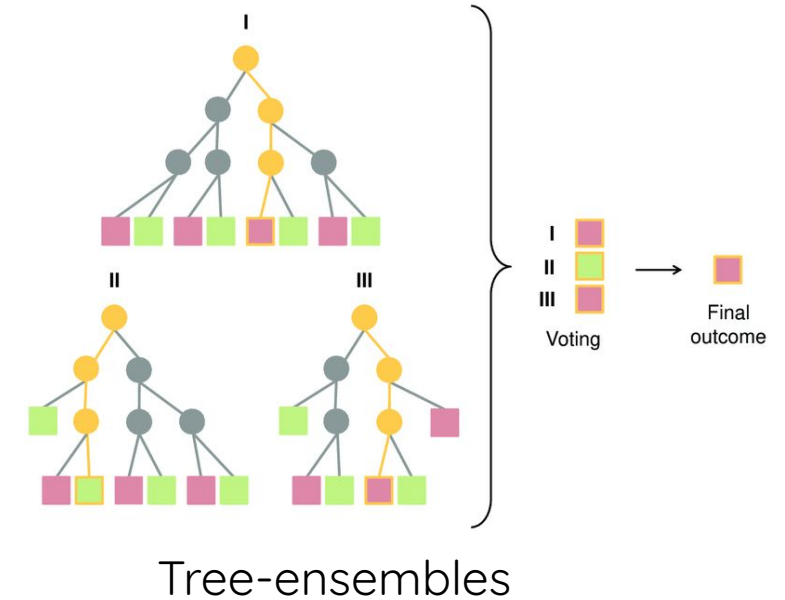
Linear Regression




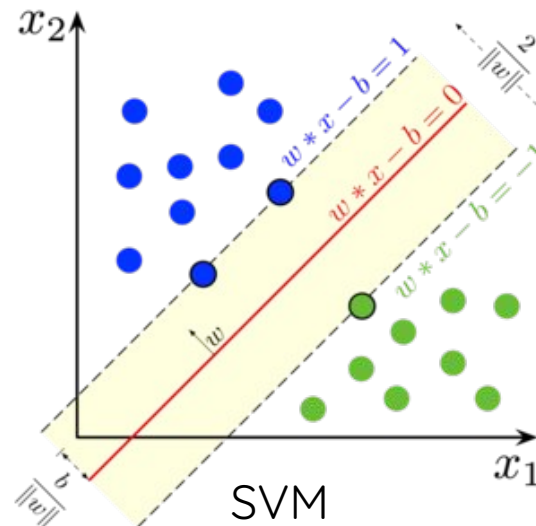
SVM

# Supervised Learning $x \rightarrow M \rightarrow y$

- Goal: *Learn* parameters (or weights) of a model ( $M$ ) that maps  $x$  to  $y$
- Example models:
  - Linear / Logistic regression
  - Support vector machines
  - Tree-ensembles: random forests, gradient boosting
  - Artificial Neural networks

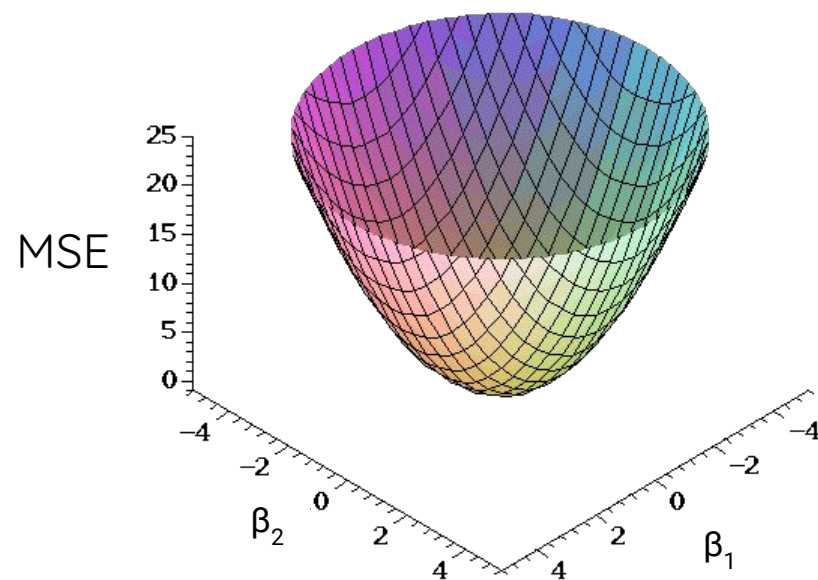
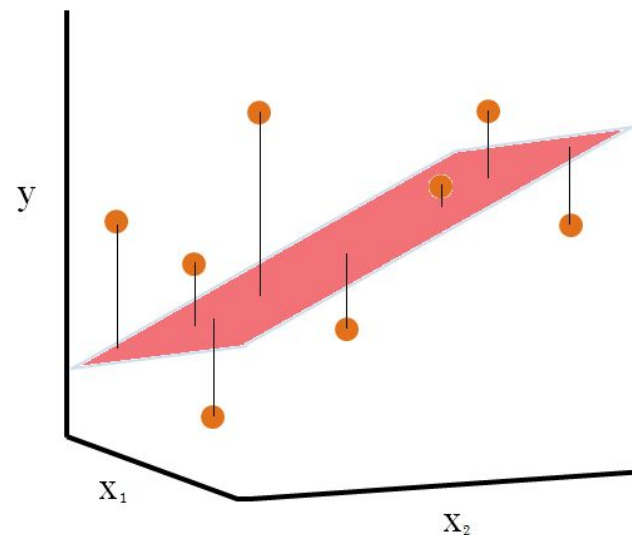


Linear Regression



# Model Fitting

- How do we learn the model weights?
  - Example: Linear regression
  - Model:  $y = \beta_0 + \beta_1 x_1^2 + \beta_2 x_2^2$
  - Loss function:  $MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$
  - Optimization: Gradient descent

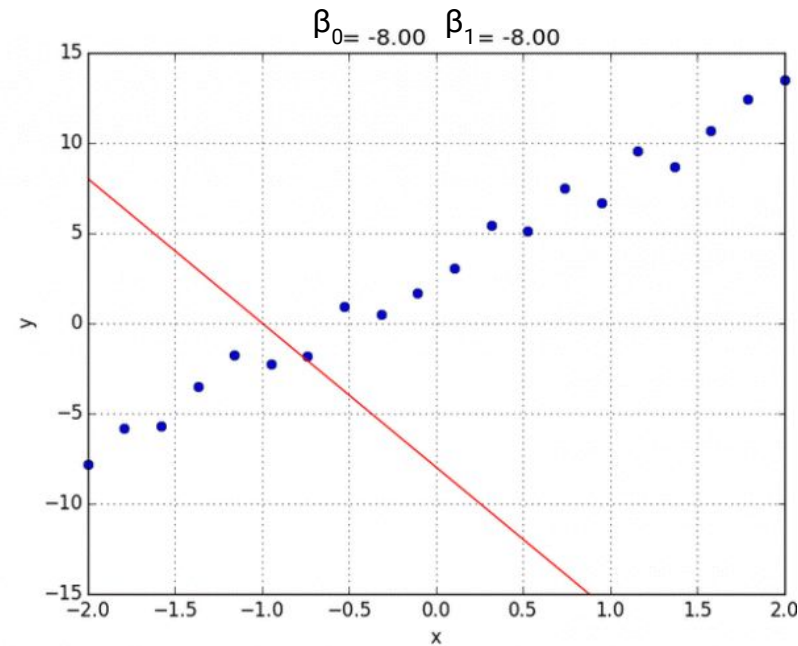
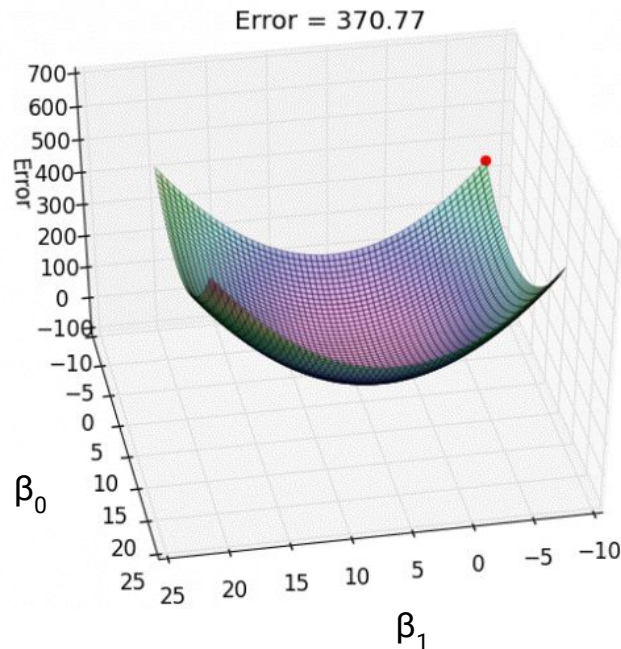


# Model Fitting

- Gradient descent with a **single** input variable and **n** samples
  - Start with random weights ( $\beta_0$  and  $\beta_1$ )
  - Compute loss (i.e. MSE)
  - Update weights based on the gradient

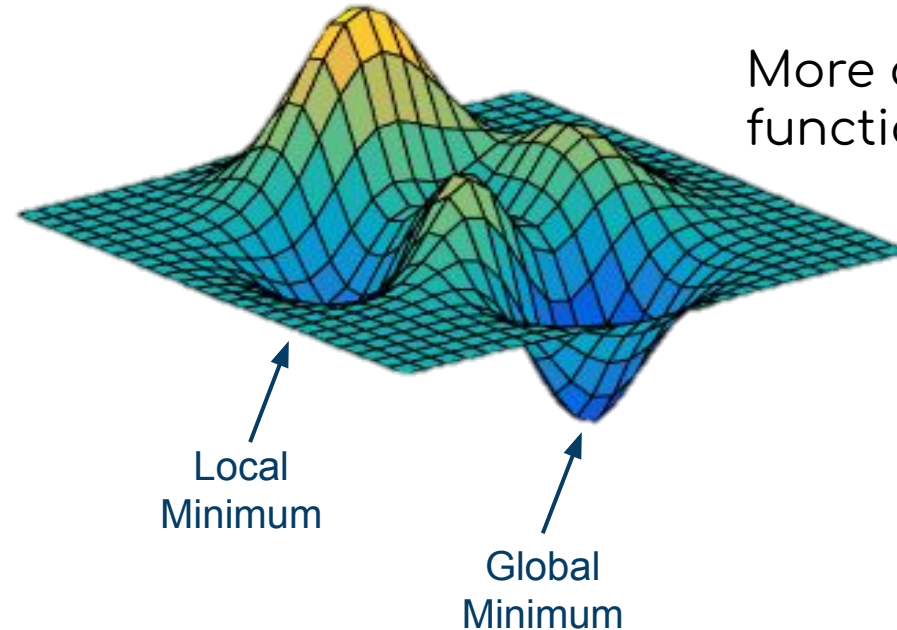
$$\hat{y}_i = \beta_0 + \beta_1 x_i$$

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$



# Model Fitting

- Gradient descent for complex models with non-convex loss functions
  - Start with random weights ( $\beta_0$  and  $\beta_1$ )
  - Compute loss
  - Update weights based on the gradient



More complex models / loss functions (e.g. ANNs)

# Model Regularization

- Why do we need to do it?
  - We have strong prior beliefs about what is a **plausible** model
    - e.g. I believe this symptom can be predicted with handful of genes.
  - Practical reasons
    - Prevent overfitting ( $n_{\text{features}} \gg n_{\text{samples}}$ )

# Model Regularization

- Why do we need to do it?
  - We have strong prior beliefs about what is a **plausible** model
    - e.g. I believe this symptom can be predicted with handful of genes.
  - Practical reasons
    - Prevent overfitting ( $n\_features \gg n\_samples$ )
- How do we do it?
  - Modify the loss function
  - Constrain the learning process

- 1) L1/Lasso: constrains parameters to be **sparse**

$$MSE = \sum_{i=1}^n (y_i - \underbrace{[\beta_0 + \sum_{j=1}^p x_{ij} \beta_j]}_{\hat{y}_i})^2 + \underbrace{\lambda \sum_{j=1}^p |\beta_j|}_{L_1}$$

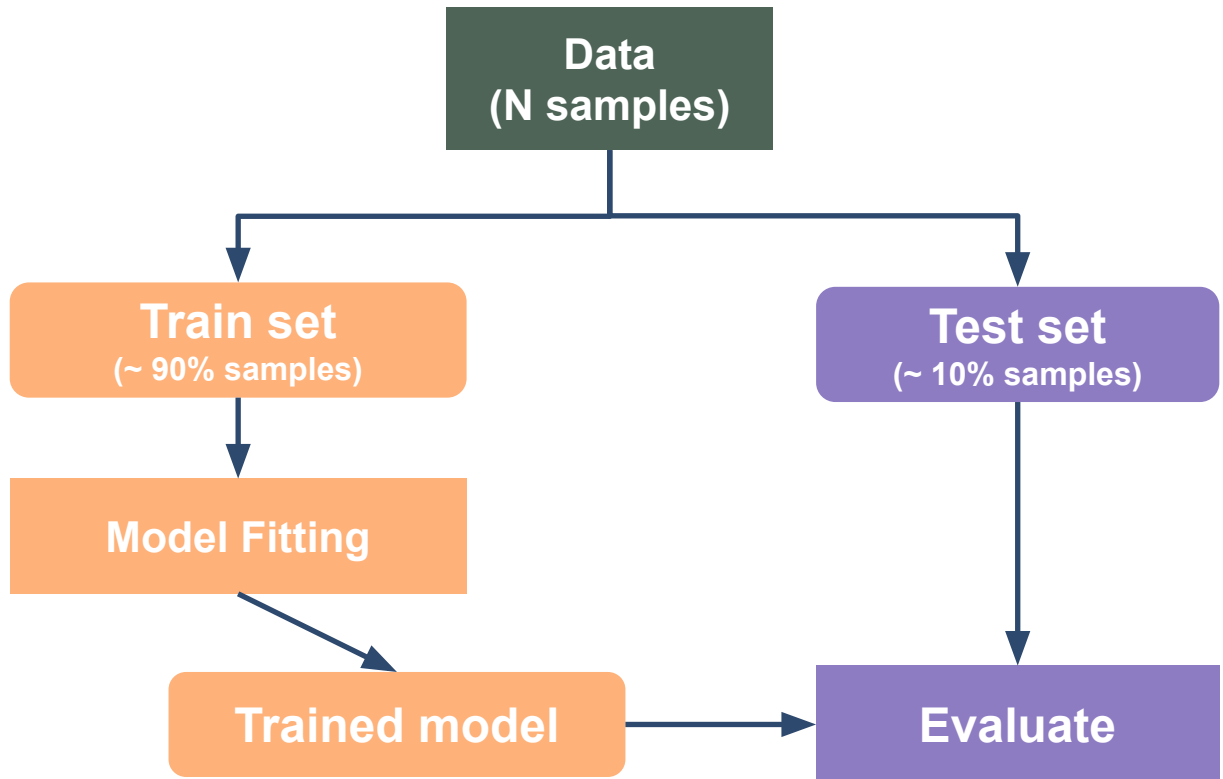
- 2) L2/Ridge: constrains parameters to be **small**

$$MSE = \sum_{i=1}^n (y_i - \underbrace{[\beta_0 + \sum_{j=1}^p x_{ij} \beta_j]}_{\hat{y}_i})^2 + \underbrace{\lambda \sum_{j=1}^p \beta_j^2}_{L_2}$$




# Model Evaluation

- What is model generalizability?
- How do we sample train and test sets?
- How do we select a model?







# Model Evaluation

- **What is model generalizability?**
- How do we sample train and test sets?
- How do we select a model?



# Model Evaluation





- Train performance  $\neq$  Test performance
  - Model: Underfitting vs Overfitting
  - Errors: Bias - Variance tradeoff
  - Regression example



● Train set

# Model Evaluation

- Train performance  $\neq$  Test performance
  - Model: Underfitting vs Overfitting
  - Errors: Bias - Variance tradeoff
  - Regression example





● Train set


● Test set

# Model Evaluation


- Train performance  $\neq$  Test performance
  - Model: Underfitting vs Overfitting
  - Errors: Bias - Variance tradeoff
  - Classification example



Underfitting



Optimal





Overfitting

○ Train class\_1


✕ Train class\_2

# Model Evaluation


- Train performance  $\neq$  Test performance
  - Model: Underfitting vs Overfitting
  - Errors: Bias - Variance tradeoff
  - Classification example



Underfitting



Optimal



Overfitting

● Train class\_1


✕ Train class\_2

● Test class\_1

✕ Test class\_2


# Model Evaluation

- What is model generalizability?
- **How do we sample train and test sets?**
- How do we select a model?



# Model Cross-Validation

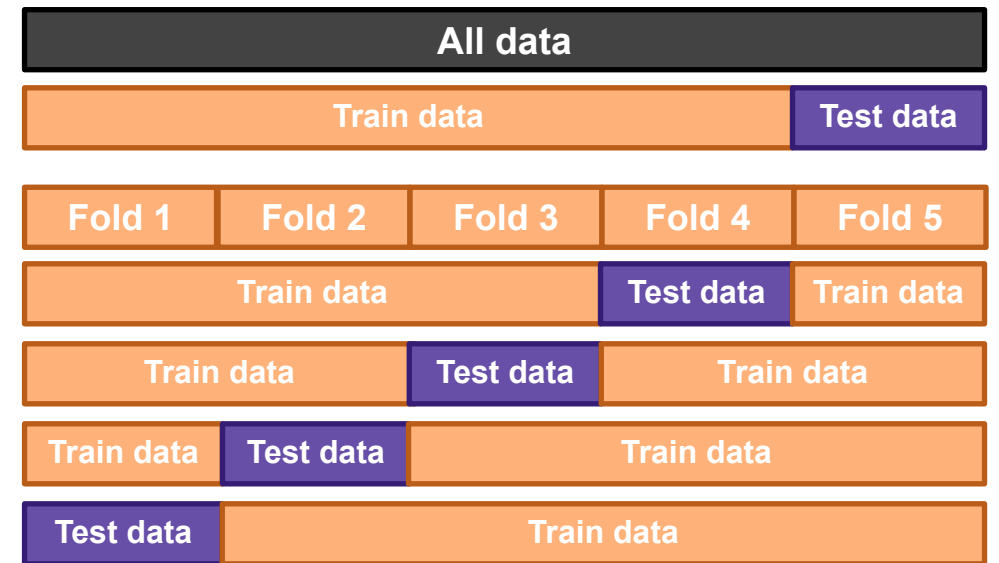
- How do we sample train and test sets?
  - Train set: learn model parameters
  - Test set (a.k.a held-out sample): Evaluate model performance





# Model Cross-Validation


- How do we sample train and test sets?
  - Train set: learn model parameters
  - Test set (a.k.a held-out sample): Evaluate model performance
  - Repeat for different Train-Test splits
    - k-fold, shuffle-split
  - Report performance statistics over all test folds



CV outer loop

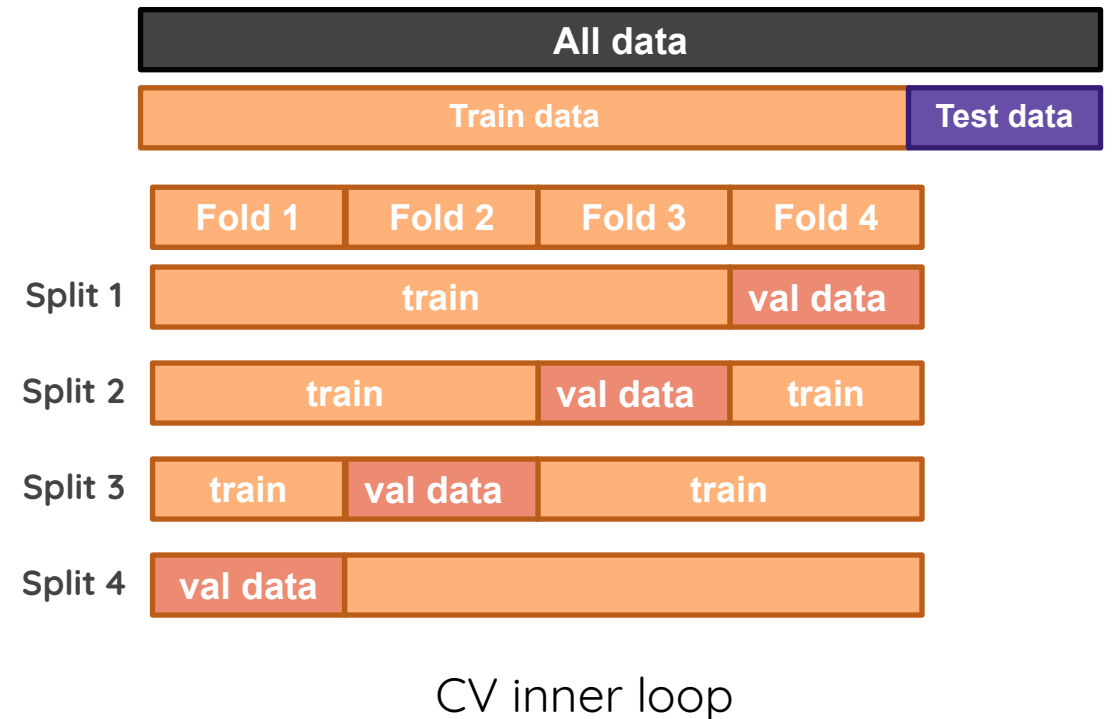
# Model Evaluation

- What is model generalizability?
- How do we sample train and test sets?
- **How do we select a model?**



# Model Cross-Validation

- How do we select a model?
  - Tune hyper-parameters of a model
  - Compare several different model architectures
  - Select / transform raw features
- This repeats for all train-test splits in the outer loop



# Hyper-parameters

- Hyper-parameter  $\neq$  parameter (or weights)
  - Parameters are **learned**; hyper-parameters are **chosen**!

# Hyper-parameters

- Hyper-parameter  $\neq$  parameter (or weights)
  - Parameters are **learned**; hyper-parameters are **chosen**!
- Examples:
  - Degree of model (eg. linear vs quadratic)
  - Kernels
  - Number of trees
  - Number of layers, filters, batch-size, learning-rate in ANNs

# Hyper-parameters

- Hyper-parameter  $\neq$  parameter (or weights)
  - Parameters are **learned**; hyper-parameters are **chosen**!
- Examples:
  - Degree of model (eg. linear vs quadratic)
  - Kernels
  - Number of trees
  - Number of layers, filters, batch-size, learning-rate in ANNs
- How do we choose them?
  - Prior beliefs  $\rightarrow$  eg. cortical thickness and age have quadratic relationship.
  - Arbitrarily  $\rightarrow$  we gotta start with something!
  - Trial and error  $\rightarrow$  do a computationally feasible grid-search.

# Hands-on Exercise 1

[https://github.com/neurodatascience/  
main-2021-ml-parts-1-2](https://github.com/neurodatascience/main-2021-ml-parts-1-2)



# Performance Scores

- Loss functions → computationally well-suited metrics
  - May / need not completely capture performance metrics of interest
- Scores → practically useful metrics
  - Binary classification

Confusion Matrix		Ground Truth	
		POSITIVE	NEGATIVE
Prediction	POSITIVE	TP	FP
	NEGATIVE	FN	TN

## False Positive



## False Negative





# Performance Scores

- ML model that detects Covid from chest CTs. Current Covid prevalence ~ 1 in 1000.
  - FP: model predicts *Covid* when person is *healthy*
  - FN: model predicts *healthy* when person has *Covid*
- What happens if we build model that predicts everyone as healthy?
  - i.e. zero FPs!



# Performance Scores

- ML model that detects Covid from chest CTs. Current Covid prevalence ~ 1 in 1000.
  - FP: model predicts *Covid* when person is *healthy*
  - FN: model predicts *healthy* when person has *Covid*
- What happens if we build model that predicts everyone as healthy?

Score	Formula	Null	What does it tell us?	When do I use it?
<b>Accuracy</b>	$(TP+TN) / (TP+FP+FN+TN)$	0.999	How many people did we correctly predict out of all the people scanned?	FNs & FPs have similar costs
<b>Precision</b>	$TP/(TP+FP)$	NaN	How many of those who we predicted as “covid” do actually have “covid”?	If you want to be more confident of your TPs
<b>Recall (aka Sensitivity)</b>	$TP/(TP+FN)$	0	Of all the people who have covid, how many of those did we correctly predict?	If you prefer FPs over FNs.
<b>Specificity</b>	$TN/(TN+FP)$	0.999	Of all the people who are healthy, how many of those did we correctly predict?	If you prefer FNs over FPs.
<b>F1</b>	$2*(Recall * Precision) / (Recall + Precision)$	NaN	Harmonic mean(average) of the precision and recall.	When you have an uneven class distribution


# Performance Curves

- Receiver Operating Characteristic (ROC) → Want high area-under-the-curve (AUC)
- Precision-Recall → Want high AUC or high Average precision (AP)



# Deep-learning


- Why the buzz?
  - Works amazing on structured input
  - Highly flexible → universal function approximator
- What are the challenges?
  - Large number of parameters → data hungry
  - Large number of hyper-parameters → difficult to train
- When do I use it?
  - If you have highly-structured input, eg. medical images.
  - You have a lot of data and computational resources.



ANN for handwritten-digit images  
(gif source: [3b1b](#))


# Pitfalls and Challenges

- Models do not generalize even after good CV performance
  - Implicit double-dipping
  - Dataset biases (eg. North-American demographics)
  - Noisy labels (eg. diagnosis definitions)
  - Data distribution shifts (eg. assay, scanner upgrades)



# Pitfalls and Challenges

- Models do not generalize even after good CV performance
  - Implicit double-dipping
  - Dataset biases (eg. North-American demographics)
  - Noisy labels (eg. diagnosis definitions)
  - Data distribution shifts (eg. assay, scanner upgrades)
- Unnecessary complexity
  - Do I really need a giant deep-net or a simple linear model would do?



# ML Novice Checklist

- Data
  - What is my n\_features and n\_samples?
  - Am I encoding categorical data correctly?
  - Am I using information (e.g. mean) from test set to preprocess (eg. zscore) the data?

# ML Novice Checklist

## ○ Data

- What is my n\_features and n\_samples?
- Am I encoding categorical data correctly?
- Am I using information (e.g. mean) from test set to preprocess (eg. zscore) the data?


## ○ Model

- Do my performance metrics capture the practical use-case of interest?
- What is the null / dummy model performance?
  - Classification: Predict majority class all the time
  - Regression: Predict the median value all the time
- Am I interpreting model coefficients correctly?



# Takeaways

- Supervised models are useful for **predictions**
  - eg. image segmentation, prognosis, drug development
- Our job is to ensure **generalizability** of these models
  - Multitude of validations
  - Understanding model biases and limitations
- Food for thought: **engineering tools** vs *scientific discovery*
  - Interpretability and explainability
  - Causality, reliability, fairness



Explainable AI

# Hands-on Exercise 2

[https://github.com/neurodatascience/  
main-2021-ml-parts-1-2](https://github.com/neurodatascience/main-2021-ml-parts-1-2)

